a Runs that reached benchmark answers during the run

Highest tier per run, from the traces (web pages themselves are not logged)

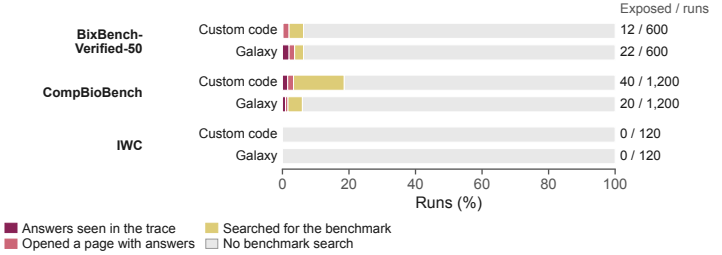

b Galaxy minus custom code without exposed runs

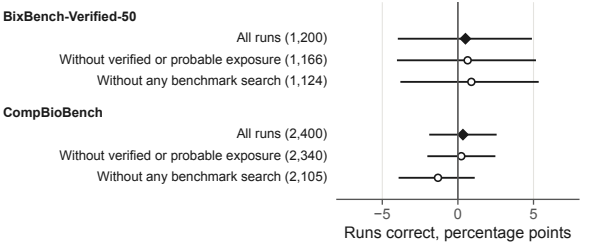

c Second rater for the failure-cause audit (BixBench-Verified-50)

45 incorrect runs; agreement 78%, Cohen's κ = 0.56; 96% counting secondary causes

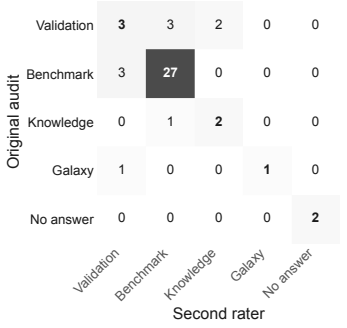

d Second rater for the failure classes of failed requests

150 requests, 10 per class; agreement 92% for classified requests (κ = 0.92); the rule's improvement group was among the rater's in 89%

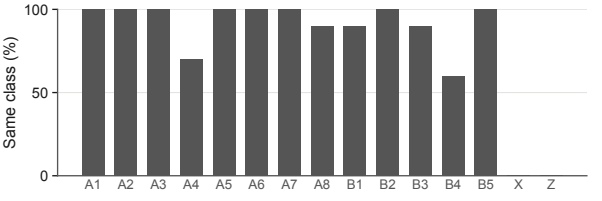

Rater judged the failure an agent error only in 3%; more than one improvement plausible in 35%