# Accuracy under three counting rules

For each model, three bars per condition: accuracy per run (correct runs ÷ runs), then per task with a task solved when ≥ 2 of its 3 runs are correct, then only when all 3 are.
Table: accuracy (%) under each rule, and Δ = Galaxy − custom code for all four models. IWC runs count as correct at ≥ 0.99 output agreement.
Paired cluster sign-flip test of Galaxy vs custom code, Holm-adjusted over each rule's 12 model comparisons: * adj. P < 0.05, ** < 0.01, *** < 0.001; no mark or ns: not significant.
Intervals and exact P values: fig2_statistical_numbers, fig2_statistical_unanimous and fig2_statistical_unanimous3.

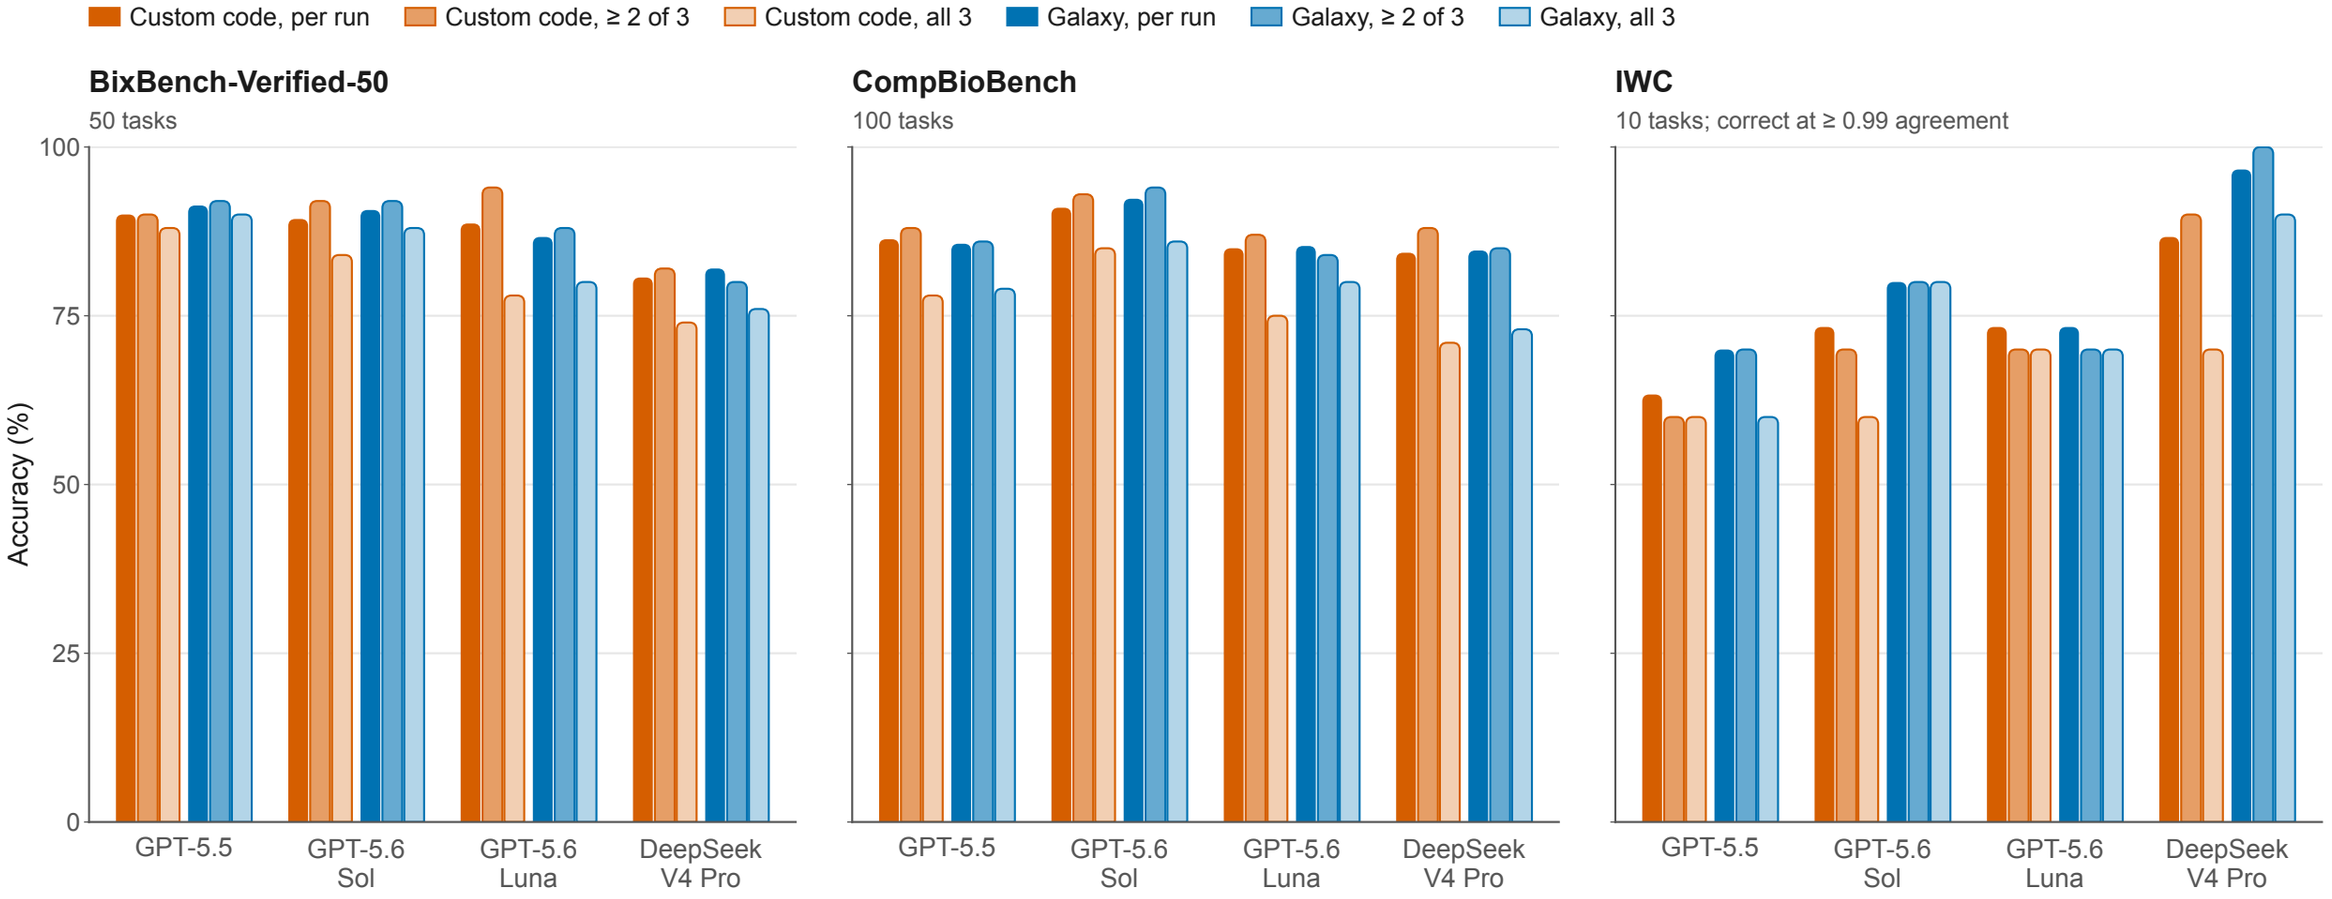


| | Per run | | ≥ 2 of 3 | | All 3 | |
|---|---|---|---|---|---|---|
| **Accuracy (%)** | Code | Galaxy | Code | Galaxy | Code | Galaxy |
| GPT-5.5 | 90.0 | 91.3 | 90.0 | 92.0 | 88.0 | 90.0 |
| GPT-5.6 Sol | 89.3 | 90.7 | 92.0 | 92.0 | 84.0 | 88.0 |
| GPT-5.6 Luna | 88.7 | 86.7 | 94.0 | 88.0 | 78.0 | 80.0 |
| DeepSeek V4 Pro | 80.7 | 82.0 | 82.0 | 80.0 | 74.0 | 76.0 |
| **All four, Δ (pts)** | **+0.5 (ns)** | | **−1.5 (ns)** | | **+2.5 (ns)** | |

| | Per run | | ≥ 2 of 3 | | All 3 | |
|---|---|---|---|---|---|---|
| **Accuracy (%)** | Code | Galaxy | Code | Galaxy | Code | Galaxy |
| GPT-5.5 | 86.3 | 85.7 | 88.0 | 86.0 | 78.0 | 79.0 |
| GPT-5.6 Sol | 91.0 | 92.3 | 93.0 | 94.0 | 85.0 | 86.0 |
| GPT-5.6 Luna | 85.0 | 85.3 | 87.0 | 84.0 | 75.0 | 80.0 |
| DeepSeek V4 Pro | 84.3 | 84.7 | 88.0 | 85.0 | 71.0 | 73.0 |
| **All four, Δ (pts)** | **+0.3 (ns)** | | **−1.8 (ns)** | | **+2.3 (ns)** | |

| | Per run | | ≥ 2 of 3 | | All 3 | |
|---|---|---|---|---|---|---|
| **Accuracy (%)** | Code | Galaxy | Code | Galaxy | Code | Galaxy |
| GPT-5.5 | 63.3 | 70.0 | 60.0 | 70.0 | 60.0 | 60.0 |
| GPT-5.6 Sol | 73.3 | 80.0 | 70.0 | 80.0 | 60.0 | 80.0 |
| GPT-5.6 Luna | 73.3 | 73.3 | 70.0 | 70.0 | 70.0 | 70.0 |
| DeepSeek V4 Pro | 86.7 | 96.7 | 90.0 | 100.0 | 70.0 | 90.0 |
| **All four, Δ (pts)** | **+5.8 (ns)** | | **+7.5 (ns)** | | **+10.0 (ns)** | |