






d IWC: most tasks reach near-perfect output agreement, and every run below 0.5 has an identified cause



e Galaxy was not the cause of most scored-incorrect BixBench-Verified-50 runs (246 runs, each execution trace examined)

A platform or tool defect was the primary or secondary cause in 22 of 111 scored-incorrect Galaxy-condition runs (20%) and in none of the 135 scored-incorrect open-ended code condition runs.



Adjudication confidence: 159 high, 55 moderate, 20 mixed, 12 unresolved