



a BixBench-Verified-50: similar acceptance

Pooled: open-ended code 85.2%, Galaxy 86.5% accepted




b CompBioBench: similar aggregate scores

Pooled mean: open-ended code 86.7, Galaxy 86.9 of 100




c IWC: equal or higher agreement with Galaxy

Pooled mean: open-ended code 0.940, Galaxy 0.980



d IWC: most tasks reach near-perfect agreement, and every score below 0.5 has an identified cause



Why scores fell below 0.5

- 1 Open-ended code, DeepSeek V4 Pro, replicate 3: hand-written multiple-testing correction was wrong
- 2 Open-ended code, GPT-5.5, replicates 1 and 3: sample names lost their capital letters
- 3 Wrong contig submitted: GPT-5.5 replicate 1 in both environments; DeepSeek V4 Pro replicate 2 in open-ended code
- 4 Galaxy, GPT-5.6 Luna, replicates 1 and 3: correct BWA-MEM output scored against the Bowtie2 reference (scoring artefact)

e Most rejected BixBench-Verified-50 answers were not caused by Galaxy (246 rejected runs, each examined call by call)

