



a No model configuration led on every benchmark: Galaxy-condition performance by model configuration



b Model configurations differed most in whether they wrote user-defined tools



c Input-token usage differed several-fold between model configurations and did not track performance

