



a BixBench-Verified-50: the Galaxy condition had more unanimous and fewer split replicate sets



All five model configurations, 250 replicate sets per condition: open-ended code 187 at 3/3, 33 split, 30 at 0/3; Galaxy 197, 28 and 25 (0/3: 5 of 50 for every model configuration).

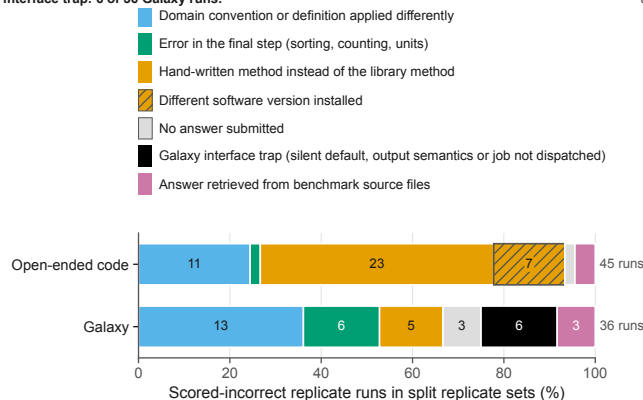
b CompBioBench: Galaxy-condition replicate sets more often returned a single distinct answer



One distinct answer in 321 of 400 open-ended code and 333 of 400 Galaxy-condition replicate sets; across all 12 runs of a condition, in 58 and 62 of 100 tasks. Answer consistency, not correctness.

c Replicate runs diverged for different reasons in the two conditions (BixBench-Verified-50)

Hand-written method or software-version difference: 30 of 45 open-ended code runs, 5 of 36 Galaxy runs;
Galaxy interface trap: 6 of 36 Galaxy runs.



d Requiring all three replicate runs to succeed widened differences between model configurations

Condition difference, Galaxy - open-ended code (percentage points):
four Codex model configurations, run-level +1.3, unanimous +3.0;
all five model configurations, run-level +3.2, unanimous +4.0.

