




a The Galaxy condition used 4–5 times more input tokens on platform-neutral benchmarks and 1.9 times more on IWC



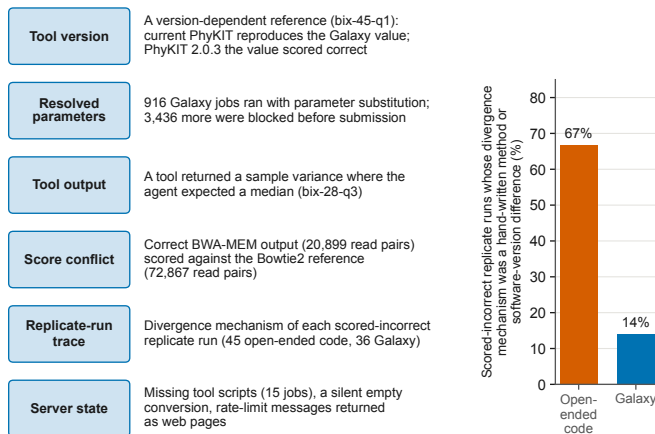
b Finding and inspecting tools took about half of the Galaxy interface activity



c Model configurations that used more input tokens were not more accurate (BixBench-Verified-50, versus GPT-5.5)



d What the Galaxy record made recoverable



Lines: 95% confidence intervals. Dashed lines: GPT-5.5 in the same execution condition.