



a Majority voting raised accuracy more in the open-ended code condition than in the Galaxy condition (BixBench-Verified-50 and CompBioBench pooled)

Change from single run: open-ended code +1.3 to +2.9 points, Galaxy -0.7 to +1.1. Majority voting changes accuracy only through split replicate sets: it gains 1/3 of a task for each set with 2 of 3 runs correct and loses 1/3 for each with 1 of 3. Open-ended code split sets mostly had 2 of 3 correct (62 versus 23); Galaxy split sets were balanced (35 versus 33), so majority voting gained little there.



b Per benchmark, majority voting raised open-ended code accuracy on CompBioBench for every model configuration; elsewhere the change depended on the model configuration



Open symbols: single run (run-level accuracy); filled symbols: majority vote; change = majority vote - single run, in points. BixBench-Verified-50: evaluator scores. CompBioBench: exact string match against the lab's score-inferred answer key (54 tasks) and score-predicted answers (46 tasks, genome-coords-q1 = E); these grades reproduce all 24 official replicate scores in the lab results file of 17 September 2026, two of which are one point above the archived values used in Fig. 2d. GPT-6 Astra, with one unpaired run per task, is excluded. IWC has a continuous endpoint, so its majority vote is the median of the three replicate runs' output agreement (for a 0/1 score, the median of three is the majority vote), over the nine tasks scored in both conditions; it is not pooled in a. Intervals: 95% percentile cluster bootstrap (source capsules for BixBench-Verified-50, tasks for CompBioBench and IWC; 20,000 resamples). Source Data give the Galaxy - open-ended code difference per benchmark and the outcome of every replicate set (majority correct, majority incorrect, no majority).