




a IWC: Galaxy runs had fewer errors than open-ended code runs (3.88 versus 4.61 per run) and recovered from 94% versus 79% of them



b BixBench-Verified-50: Galaxy runs had about as many errors as open-ended code runs (2.31 versus 2.26 per run) and recovered from 78% versus 75% of them



c CompBioBench: Galaxy runs had more errors than open-ended code runs (4.45 versus 2.87 per run) and recovered from 80% versus 75% of them



■ Open-ended code condition ■ Galaxy condition

Error: a shell command that exited with a non-zero code (a silent exit code 1, such as a search with no match, is not counted) or a Galaxy analysis job that ended in the error state. Its type comes from its error message (ordered rules in fig_on_demand.ERROR_RULES; Galaxy job messages from the archived job records); "Galaxy job never started" is a job with neither a command line nor output. Recovered: the error occurred in a run that still ended correct (BixBench-Verified-50 scored correct; CompBioBench exact match to the reference key; IWC output agreement 0.95 or above, nine tasks). Bars: errors per run; points: share recovered with 95% cluster-bootstrap interval (source capsules for BixBench-Verified-50, tasks otherwise; 20,000 resamples), shown when at least 10 errors. Excluded: 36 BixBench-Verified-50 (no detailed analysis history); 2 CompBioBench (no detailed analysis history); 12 CompBioBench (no execution trace); 2 IWC (no detailed analysis history); GPT-6 Astra.