



a Paired execution conditions (assigned arms)



b Analysed populations



c Benchmark-specific primary endpoints (four Codex model configurations)

Benchmark	Endpoint	Reference provenance	Open-ended code	Galaxy	Galaxy – open-ended code (95% CI)	Discordant replicate sets, code vs Galaxy
BixBench-Verified-50	Accepted by the original evaluator (% of runs)	Benchmark reference answer and its original evaluator	85.2	86.5	+1.3 (–3.5 to 6.3)	19 of 200 vs 15 of 200
CompBioBench	Agreement with the reconstructed reference key (% of runs)	Key reconstructed from leaderboard totals (54 items score-inferred, 46 score-predicted)	86.7	87.1	+0.4 (–1.6 to 2.4)	66 of 400 vs 53 of 400
IWC	Mean output agreement with the workflow output (0-1), nine matched tasks	Published workflow output; some route references calibrated from agent runs	0.940	0.980	+0.040 (0.010 to 0.078)	11 of 36 vs 4 of 36

BixBench-Verified-50 and IWC intervals are the archived 95% cluster-bootstrap intervals; the CompBioBench interval is from this analysis (20,000 resamples). IWC discordant sets: three output-agreement values spanning more than 0.05.