

Accuracy under three counting rules

For each model, three bars per condition: accuracy per run (correct runs ÷ runs), then per task with a task solved when ≥ 2 of its 3 runs are correct, then only when all 3 are.
Table: accuracy (%) under each rule, and Δ = Galaxy – custom code for all four models. IWC runs count as correct at ≥ 0.99 output agreement.
Paired cluster sign-flip test of Galaxy vs custom code, Holm-adjusted over each rule's 12 model comparisons: * adj. P < 0.05, ** < 0.01, *** < 0.001; no mark or ns: not significant.
Intervals and exact P values: fig2_statistical_numbers, fig2_statistical_unanimous and fig2_statistical_unanimous3.

Custom code, per run Custom code, ≥ 2 of 3 Custom code, all 3 Galaxy, per run Galaxy, ≥ 2 of 3 Galaxy, all 3

