Accuracy by model, with statistical tests of the condition difference

Custom code | Galaxy | 95% CI

Bars: accuracy with 95% cluster-bootstrap intervals (20,000 resamples; clusters are BixBench source capsules, otherwise tasks). Brackets: Galaxy vs custom code for each model, paired cluster sign-flip test (200,000 draws; exact for IWC), Holm-adjusted over the 12 model comparisons: * adj. P < 0.05, ** < 0.01, *** < 0.001, ns not significant. IWC runs count as correct at ≥ 0.99.

BixBench-Verified-50

50 tasks

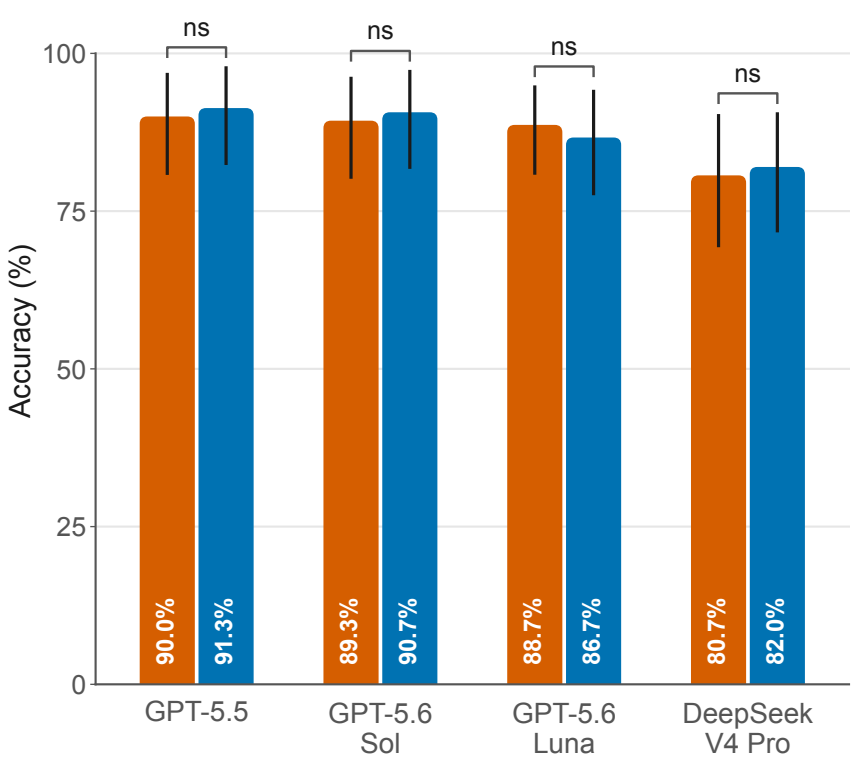

All four models: Galaxy − custom code +0.5 pts
(95% CI −3.9 to 4.8; P = 0.88, adj. over 3 benchmarks P = 1.00)

Galaxy − custom code	Δ (pts)	P	adj. P
GPT-5.5	+1.3	0.50	1.00
GPT-5.6 Sol	+1.3	0.86	1.00
GPT-5.6 Luna	−2.0	0.69	1.00
DeepSeek V4 Pro	+1.3	0.84	1.00

CompBioBench

100 tasks

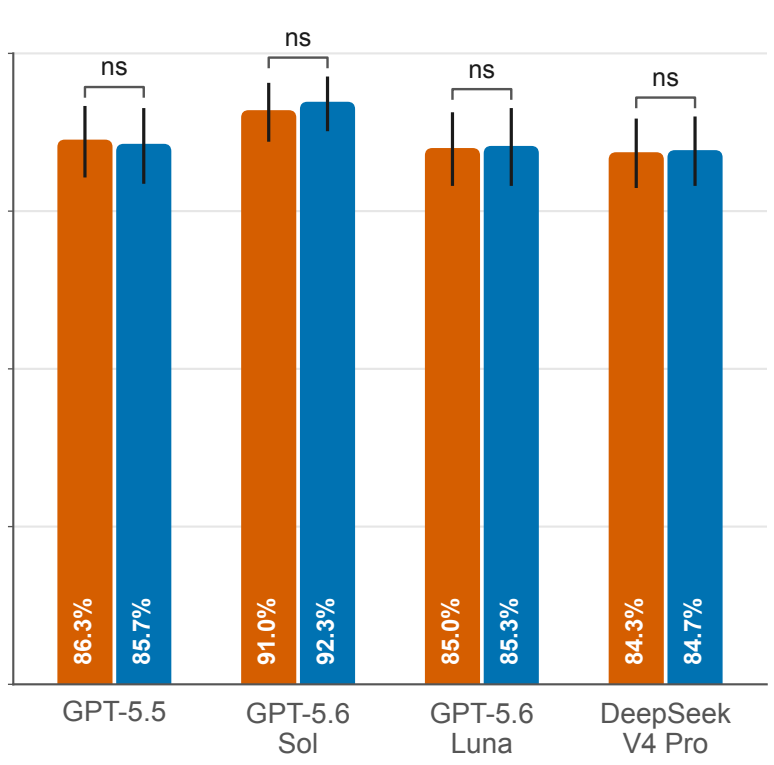

All four models: Galaxy − custom code +0.3 pts
(95% CI −1.8 to 2.6; P = 0.82, adj. over 3 benchmarks P = 1.00)

Galaxy − custom code	Δ (pts)	P	adj. P
GPT-5.5	−0.7	0.85	1.00
GPT-5.6 Sol	+1.3	0.53	1.00
GPT-5.6 Luna	+0.3	1.00	1.00
DeepSeek V4 Pro	+0.3	1.00	1.00

IWC

10 tasks; correct at ≥ 0.99 agreement

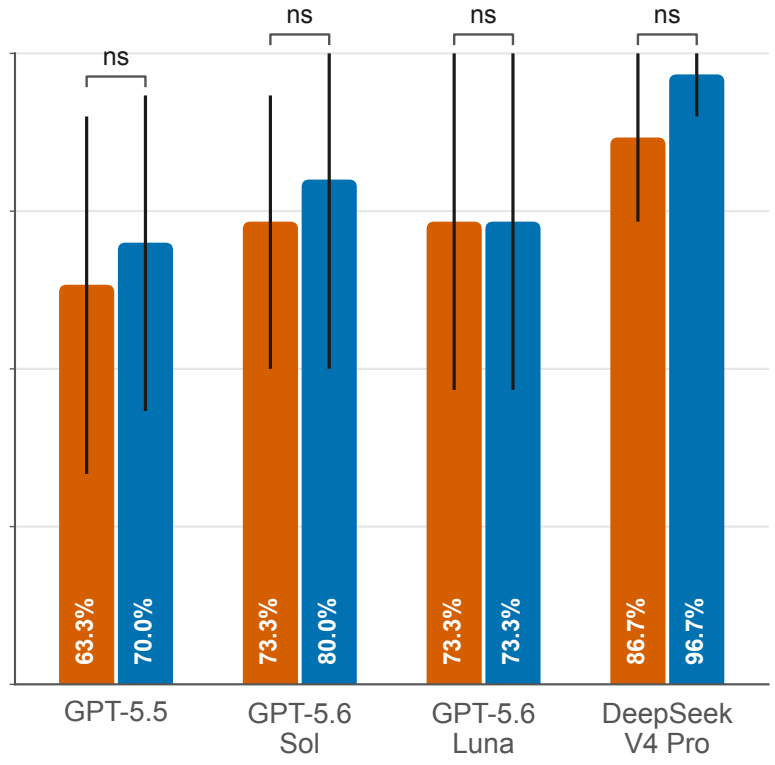

All four models: Galaxy − custom code +5.8 pts
(95% CI −0.8 to 15.0; P = 0.38, adj. over 3 benchmarks P = 1.00)

Galaxy − custom code	Δ (pts)	P	adj. P
GPT-5.5	+6.7	0.50	1.00
GPT-5.6 Sol	+6.7	0.75	1.00
GPT-5.6 Luna	+0.0	1.00	1.00
DeepSeek V4 Pro	+10.0	0.25	1.00