Task accuracy by model (task solved when ≥ 2 of 3 runs are correct)

Custom code Galaxy 95% CI

Accuracy = tasks solved ÷ tasks (BixBench-Verified-50: 50; CompBioBench: 100; IWC: 10); a model solves a task when at least 2 of its 3 runs are correct (IWC runs: ≥ 0.99 agreement).
Bars: 95% cluster-bootstrap intervals (20,000 resamples; clusters are BixBench source capsules, otherwise tasks). Brackets: Galaxy vs custom code for each model,
paired cluster sign-flip test on tasks solved (200,000 draws; exact for IWC), Holm-adjusted over the 12 model comparisons: * adj. P < 0.05, ** < 0.01, *** < 0.001, ns not significant.

BixBench-Verified-50

50 tasks

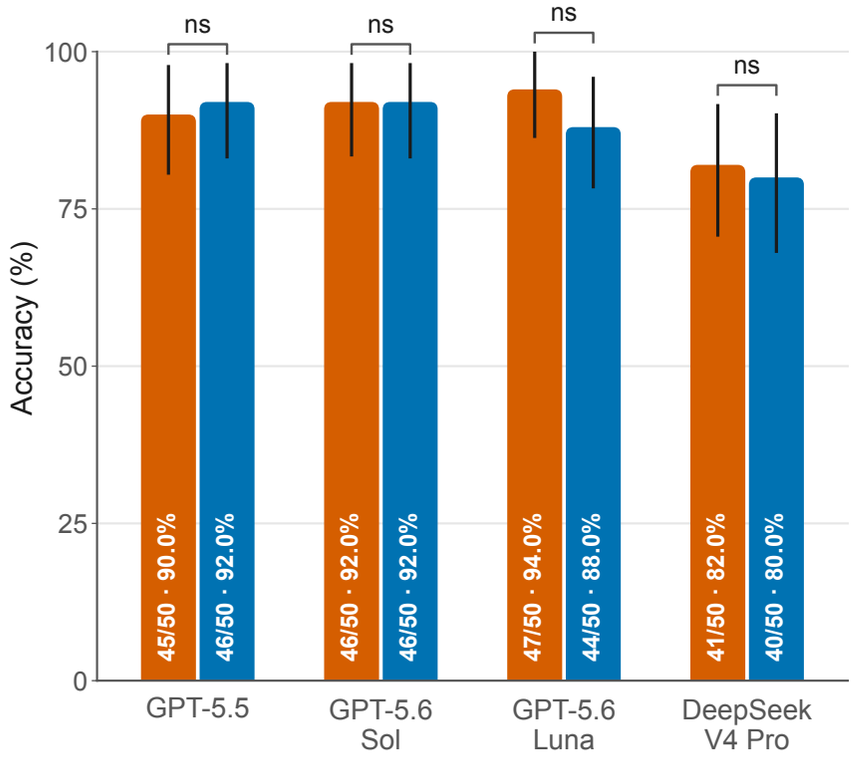

All four models: Galaxy − custom code −1.5 pts
(95% CI −5.7 to 2.0; P = 0.63, adj. over 3 benchmarks P = 1.00)

Galaxy − custom code	Δ (pts)	P	adj. P
GPT-5.5	+2.0	1.00	1.00
GPT-5.6 Sol	+0.0	1.00	1.00
GPT-5.6 Luna	−6.0	0.25	1.00
DeepSeek V4 Pro	−2.0	1.00	1.00

CompBioBench

100 tasks

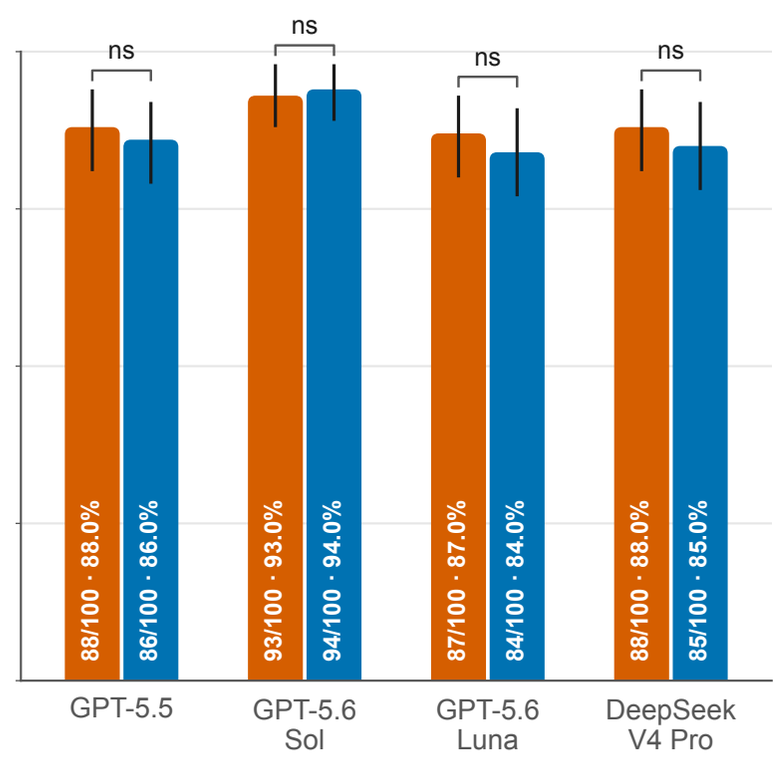

All four models: Galaxy − custom code −1.8 pts
(95% CI −4.8 to 1.0; P = 0.32, adj. over 3 benchmarks P = 0.95)

Galaxy − custom code	Δ (pts)	P	adj. P
GPT-5.5	−2.0	0.69	1.00
GPT-5.6 Sol	+1.0	1.00	1.00
GPT-5.6 Luna	−3.0	0.55	1.00
DeepSeek V4 Pro	−3.0	0.51	1.00

IWC

10 tasks; correct at ≥ 0.99 agreement

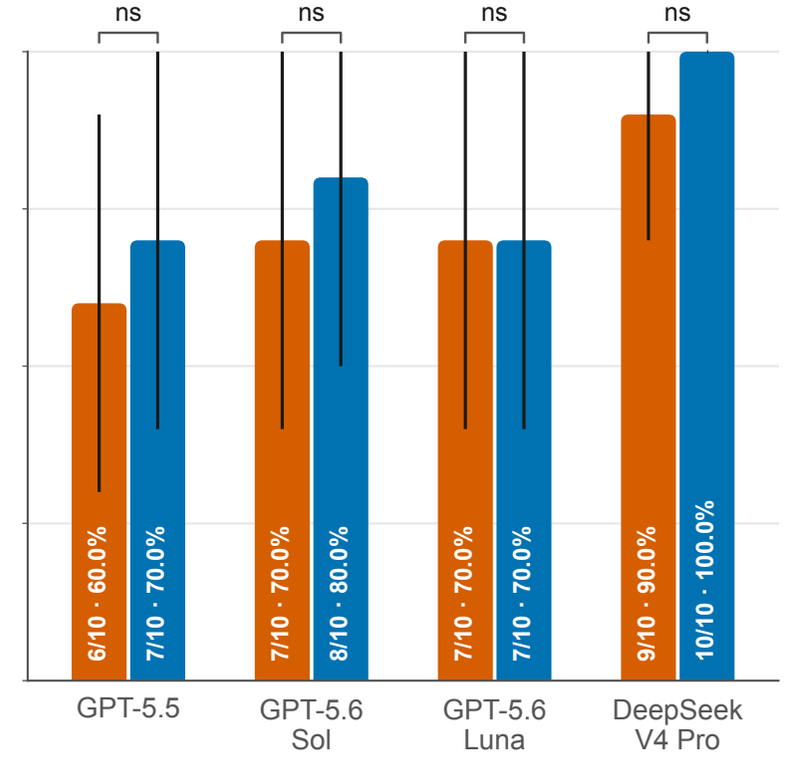

All four models: Galaxy − custom code +7.5 pts
(95% CI 0.0 to 17.5; P = 0.50, adj. over 3 benchmarks P = 1.00)

Galaxy − custom code	Δ (pts)	P	adj. P
GPT-5.5	+10.0	1.00	1.00
GPT-5.6 Sol	+10.0	1.00	1.00
GPT-5.6 Luna	+0.0	1.00	1.00
DeepSeek V4 Pro	+10.0	1.00	1.00