

Runs correct per task: custom code against Galaxy

Same number correct Galaxy more correct Custom code more correct

Each cell counts the tasks of one model by how many of their 3 runs were correct with custom code (x) and with Galaxy (y); the conditions are paired by task. Diagonal: the same number correct in both; above it, Galaxy got more runs right; below it, custom code did. Shade: number of tasks (log scale). Dashed lines: the ≥ 2-of-3 rule of fig2_task_similarity (bottom left: both fail; top left: only custom code fails; bottom right: only Galaxy fails; top right: neither fails). IWC runs count as correct at ≥ 0.99 output agreement.

BixBench-Verified-50

50 tasks


CompBioBench

100 tasks

IWC

10 tasks; correct at ≥ 0.99 agreement


GPT-5.5




GPT-5.5




GPT-5.5




GPT-5.6 Sol




GPT-5.6 Sol




GPT-5.6 Sol




GPT-5.6 Luna




GPT-5.6 Luna




GPT-5.6 Luna



DeepSeek V4 Pro



DeepSeek V4 Pro



DeepSeek V4 Pro

