More tokens do not buy accuracy

a, accuracy against median tokens per run for each model and condition; a grey line joins a model's two conditions. Galaxy always costs more tokens, and its accuracy is higher for some models and lower for others. b, the same task, model and condition: tokens of the incorrect runs relative to the correct ones, the geometric mean over the replicate sets that had both outcomes (n sets; at least 3 needed), with a 95% cluster-bootstrap interval. Above 1, failing runs used more tokens. IWC: correct at ≥ 0.99.

Custom code Galaxy GPT-5.5 GPT-5.6 Sol GPT-5.6 Luna DeepSeek V4 Pro

a

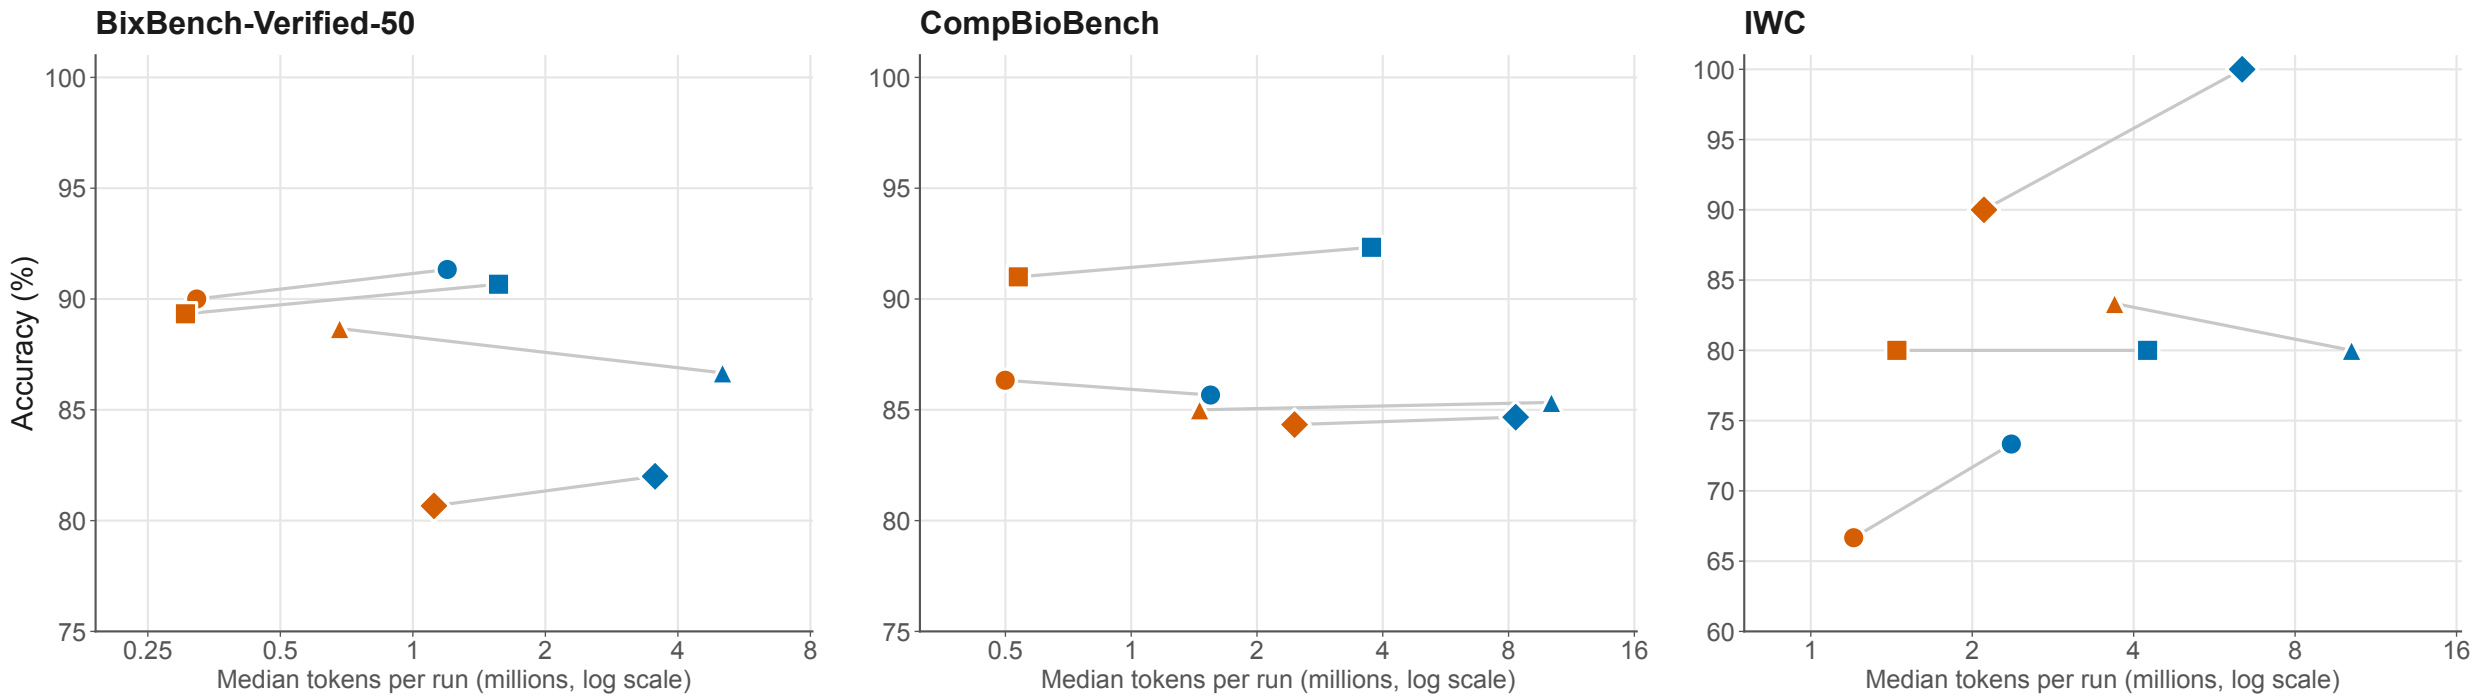
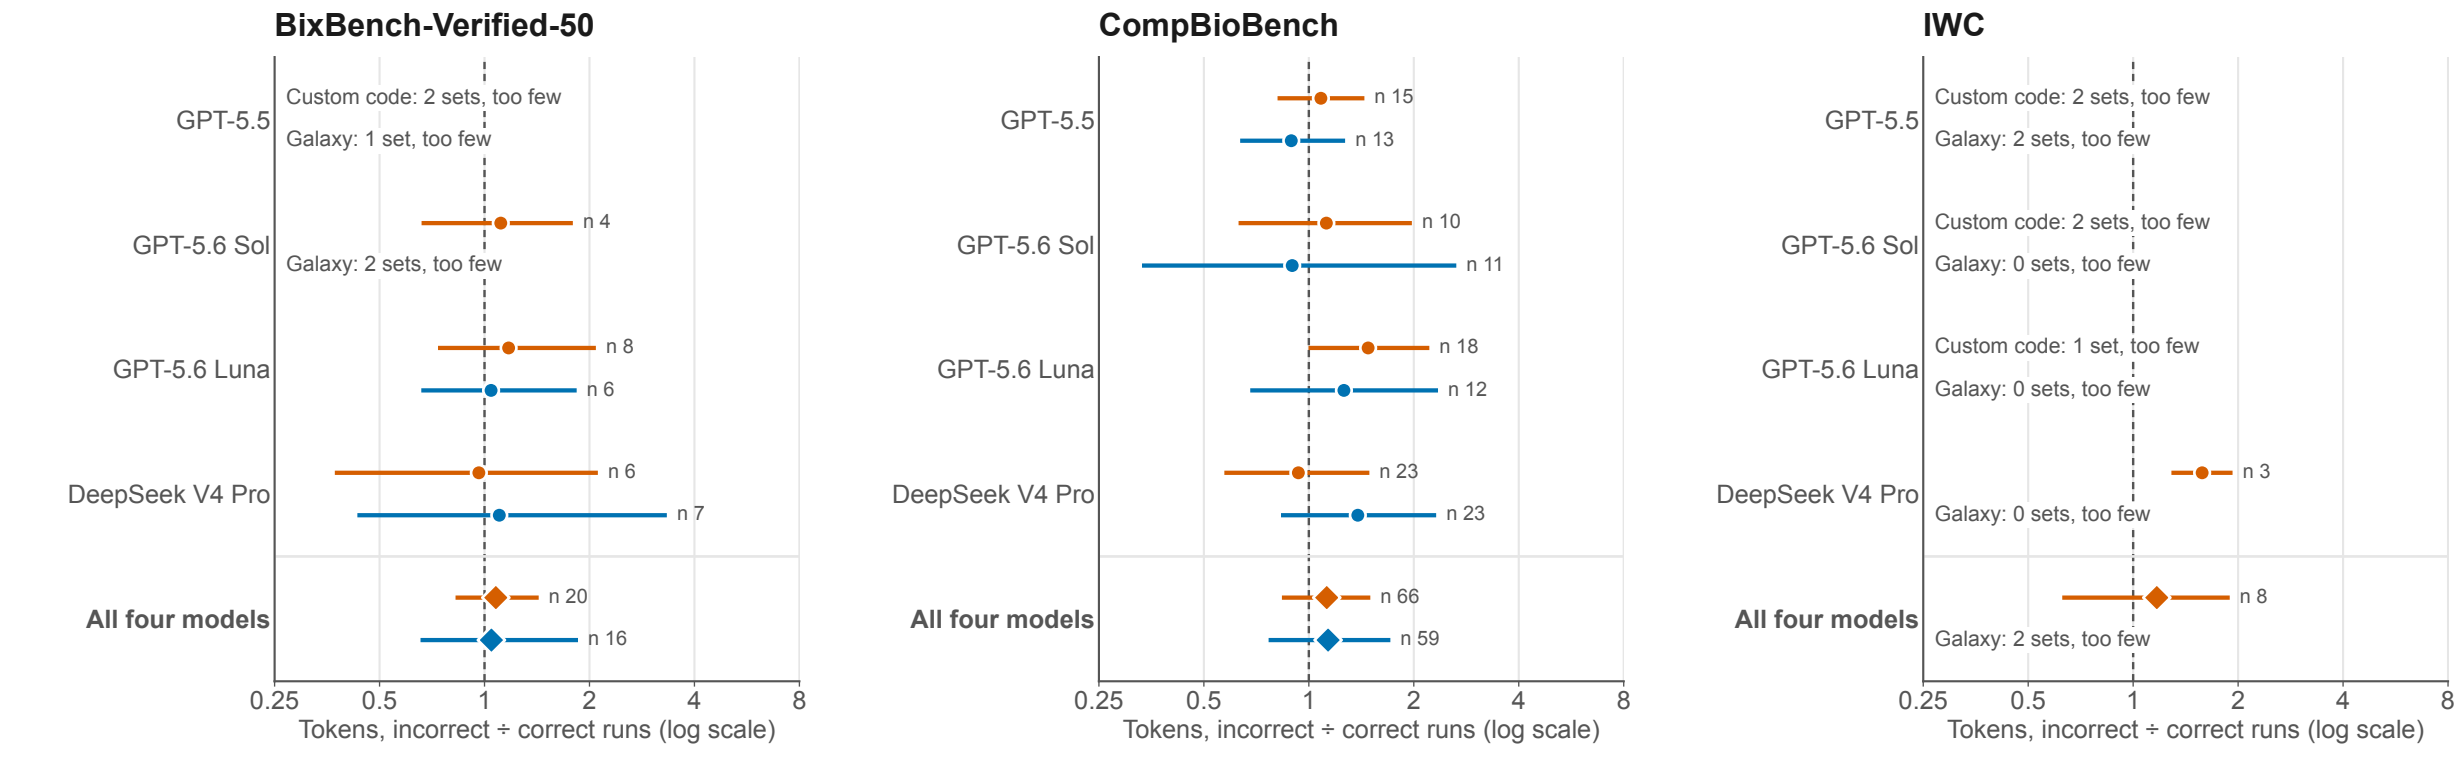
b