


Interface changes cut Galaxy tokens by more than half, with accuracy unchanged

BixBench-Verified-50. a, b, tokens per complete 50-task run (input including cached context, plus output; mean of 3 runs per task), Galaxy and custom code side by side at each stage; above each pair: Galaxy ÷ custom code; in the bars: accuracy. c, each task's Galaxy tokens after the changes relative to before (3 runs each, summed); the bar is the median. d, where the October savings came from: mean input tokens per run by phase (estimated as in the phase figure). Intervals: 95% cluster bootstrap over source capsules.

Custom code Galaxy before the changes Galaxy, intermediate round Galaxy after the changes



Round 3 (1 replicate): accuracy not reported. Custom code: the same archived runs at each stage.

